

```
In [22]: import pandas as pd
import numpy as np
import sklearn
import matplotlib.pyplot as plt
```

```
In [23]: data = pd.read_csv('Thyroid_Diff.csv')
```

```
In [24]: data.head()
```

```
Out[24]:
```

| | Age | Gender | Smoking | Hx Smoking | Hx Radiothreapy | Thyroid Function | Physical Examination | Adenopathy | Patholog |
|---|-----|--------|---------|---------------|--------------------|---------------------|-----------------------------------|------------|---------------|
| 0 | 27 | F | No | No | No | Euthyroid | Single nodular goiter-left | No | Micropapillar |
| 1 | 34 | F | No | Yes | No | Euthyroid | Multinodular goiter | No | Micropapillar |
| 2 | 30 | F | No | No | No | Euthyroid | Single nodular goiter-right | No | Micropapillar |
| 3 | 62 | F | No | No | No | Euthyroid | Single nodular goiter-right | No | Micropapillar |
| 4 | 62 | F | No | No | No | Euthyroid | Multinodular goiter | No | Micropapillar |

Data Preprocessing

Data Cleaning

```
In [25]: data.info()
#No null values found
#Preprocess object into integer/float
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 383 entries, 0 to 382
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    383 non-null    int64
1   Gender                 383 non-null    object
2   Smoking                383 non-null    object
3   Hx Smoking             383 non-null    object
4   Hx Radiothreapy        383 non-null    object
5   Thyroid Function       383 non-null    object
6   Physical Examination   383 non-null    object
7   Adenopathy             383 non-null    object
8   Pathology              383 non-null    object
9   Focality               383 non-null    object
10  Risk                   383 non-null    object
11  T                       383 non-null    object
12  N                       383 non-null    object
13  M                       383 non-null    object
14  Stage                  383 non-null    object
15  Response                383 non-null    object
16  Recurred                383 non-null    object
dtypes: int64(1), object(16)
memory usage: 51.0+ KB
```

```
In [26]: data.isnull().values.any()
```

```
Out[26]: False
```


```
In [27]: data.describe().T
```

```
Out[27]:
```

| | count | mean | std | min | 25% | 50% | 75% | max |
|------------|-------|-----------|-----------|------|------|------|------|------|
| Age | 383.0 | 40.866841 | 15.134494 | 15.0 | 29.0 | 37.0 | 51.0 | 82.0 |

```
In [28]: # extra code - the next 5 lines define the default font sizes
plt.rc('font', size=14)
plt.rc('axes', labelsz=14, titlesz=14)
plt.rc('legend', fontsize=14)
plt.rc('xtick', labelsz=10)
plt.rc('ytick', labelsz=10)

data.hist(bins=50, figsize=(12,8))
plt.show()
```



```
In [29]: # Replace values
data['Recurred'] = data['Recurred'].replace({'Yes': '1', 'No': '0'})

# Write the modified DataFrame back to CSV
data.to_csv('modified_thyroid.csv', index=False)
```

```
In [30]: df = pd.read_csv('modified_thyroid.csv')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 383 entries, 0 to 382
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    383 non-null   int64
1   Gender                 383 non-null   object
2   Smoking                383 non-null   object
3   Hx Smoking             383 non-null   object
4   Hx Radiothreapy        383 non-null   object
5   Thyroid Function       383 non-null   object
6   Physical Examination    383 non-null   object
7   Adenopathy             383 non-null   object
8   Pathology              383 non-null   object
9   Focality               383 non-null   object
10  Risk                   383 non-null   object
11  T                       383 non-null   object
12  N                       383 non-null   object
13  M                       383 non-null   object
14  Stage                  383 non-null   object
15  Response               383 non-null   object
16  Recurred               383 non-null   int64
dtypes: int64(2), object(15)
memory usage: 51.0+ KB
```

```
In [31]: recurrence= df.drop('Recurred', axis=1)
recurrence_labels = df[['Recurred']].copy()
```

```
In [32]: from sklearn.impute import SimpleImputer
imputer = SimpleImputer(strategy='median') #impute using the median value
```

```
In [33]: #Select the numerical variables
recurrence_num = recurrence.select_dtypes(include=[np.number])
```

```
In [34]: imputer.fit(recurrence_num)
```

```
Out[34]: SimpleImputer
SimpleImputer(strategy='median')
```

```
In [35]: #show the median values for each variable
imputer.statistics_
```

```
Out[35]: array([37.])
```

```
In [36]: X = imputer.transform(recurrence_num)
```

```
In [37]: imputer.feature_names_in_
```

```
Out[37]: array(['Age'], dtype=object)
```

```
In [38]: from sklearn import set_config

set_config(transform_output='pandas') #scikit-learn >= 1.2
```

```
In [39]: #Dropping outliers
from sklearn.ensemble import IsolationForest

isolation_forest = IsolationForest(random_state=42)
outlier_pred = isolation_forest.fit_predict(X)
```

```
In [40]: pd.Series(outlier_pred).value_counts()
```

```
Out[40]: 1    232
        -1   151
        Name: count, dtype: int64
```

Encode Features

```
In [41]: recurrence_cat = recurrence[['Gender', 'Smoking', 'Hx Smoking', 'Hx Radiothreapy', 'Thyroid Function',
                                     'Physical Examination', 'Adenopathy', 'Pathology', 'Focal Lesion',
                                     'Stage', 'Response']]
recurrence_cat.head()
```

```
Out[41]:
```

| | Gender | Smoking | Hx Smoking | Hx Radiothreapy | Thyroid Function | Physical Examination | Adenopathy | Pathology | Focal Lesion |
|---|--------|---------|------------|-----------------|------------------|-----------------------------|------------|----------------|--------------|
| 0 | F | No | No | No | Euthyroid | Single nodular goiter-left | No | Micropapillary | |
| 1 | F | No | Yes | No | Euthyroid | Multinodular goiter | No | Micropapillary | |
| 2 | F | No | No | No | Euthyroid | Single nodular goiter-right | No | Micropapillary | |
| 3 | F | No | No | No | Euthyroid | Single nodular goiter-right | No | Micropapillary | |
| 4 | F | No | No | No | Euthyroid | Multinodular goiter | No | Micropapillary | M |

```
In [42]: recurrence_cat_oneHot = recurrence[['Thyroid Function', 'Physical Examination', 'Pathology', 'Focal Lesion',
                                             'Stage', 'Response']]
recurrence_cat_ordinal = recurrence[['Gender', 'Smoking', 'Hx Smoking', 'Hx Radiothreapy', 'Adenopathy', 'Pathology', 'Focal Lesion',
                                     'Stage', 'Response']]
```


```
In [43]: from sklearn.preprocessing import OneHotEncoder, LabelEncoder

cat_encoder = OneHotEncoder(sparse_output=False, handle_unknown='ignore')
recurrence_cat_1hot = cat_encoder.fit_transform(recurrence_cat_oneHot)
recurrence_cat_1hot
```

Out[43]:

| | Thyroid Function_Clinical Hyperthyroidism | Thyroid Function_Clinical Hypothyroidism | Thyroid Function_Euthyroid | Thyroid Function_Subclinical Hyperthyroidism | Thyroid Function_Subclinical Hypothyroidism |
|-----|---|--|-------------------------------|--|---|
| 0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| 1 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| 2 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| 3 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| 4 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| ... | ... | ... | ... | ... | ... |
| 378 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| 379 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| 380 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |
| 381 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 382 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 |

383 rows × 46 columns



In [44]:

```
recurrence_cat_ordinal
```

Out[44]:

| | Gender | Smoking | Hx Smoking | Hx Radiothreapy |
|-----|--------|---------|------------|-----------------|
| 0 | F | No | No | No |
| 1 | F | No | Yes | No |
| 2 | F | No | No | No |
| 3 | F | No | No | No |
| 4 | F | No | No | No |
| ... | ... | ... | ... | ... |
| 378 | M | Yes | Yes | Yes |
| 379 | M | Yes | No | Yes |
| 380 | M | Yes | Yes | No |
| 381 | M | Yes | Yes | Yes |
| 382 | M | Yes | No | No |

383 rows × 4 columns

Encode Target

In [45]:

```
recurrence_tar_ordinal = df[['Recurred']]
```

In [46]: `recurrence_tar_ordinal`

Out[46]:

| | Recurred |
|-----|----------|
| 0 | 0 |
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 0 |
| ... | ... |
| 378 | 1 |
| 379 | 1 |
| 380 | 1 |
| 381 | 1 |
| 382 | 1 |

383 rows × 1 columns

In [47]: `cat_encoder.categories_`

Out[47]:

```
[array(['Clinical Hyperthyroidism', 'Clinical Hypothyroidism', 'Euthyroid',
        'Subclinical Hyperthyroidism', 'Subclinical Hypothyroidism'],
      dtype=object),
 array(['Diffuse goiter', 'Multinodular goiter', 'Normal',
        'Single nodular goiter-left', 'Single nodular goiter-right'],
      dtype=object),
 array(['Follicular', 'Hurthel cell', 'Micropapillary', 'Papillary'],
      dtype=object),
 array(['Multi-Focal', 'Uni-Focal'], dtype=object),
 array(['High', 'Intermediate', 'Low'], dtype=object),
 array(['T1a', 'T1b', 'T2', 'T3a', 'T3b', 'T4a', 'T4b'], dtype=object),
 array(['N0', 'N1a', 'N1b'], dtype=object),
 array(['M0', 'M1'], dtype=object),
 array(['I', 'II', 'III', 'IVA', 'IVB'], dtype=object),
 array(['Biochemical Incomplete', 'Excellent', 'Indeterminate',
        'Structural Incomplete'], dtype=object),
 array(['Bilateral', 'Extensive', 'Left', 'No', 'Posterior', 'Right'],
      dtype=object)]
```

Feature Scaling

In [48]:

```
from sklearn.preprocessing import StandardScaler
std_scaler = StandardScaler()
recurrence_scaled = std_scaler.fit_transform(recurrence_num)
```

In [49]: `recurrence_num.columns`


Out[49]: Index(['Age'], dtype='object')

Transformation Pipeline

```
In [50]: from sklearn.pipeline import Pipeline

num_pipeline = Pipeline([
    ("impute", SimpleImputer(strategy='median')),
    ('standardize', StandardScaler())
])
```

```
In [51]: set_config(display='diagram')
num_pipeline
```



```
In [52]: recurrence_prepared = num_pipeline.fit_transform(recurrence_num)
recurrence_prepared.head()
```

Out[52]:

| | Age |
|---|-----------|
| 0 | -0.917439 |
| 1 | -0.454315 |
| 2 | -0.718957 |
| 3 | 1.398184 |
| 4 | 1.398184 |

```
In [53]: recurrence_prepared.describe()
```

Out[53]:

| | Age |
|-------|---------------|
| count | 3.830000e+02 |
| mean | -1.855203e-17 |
| std | 1.001308e+00 |
| min | -1.711367e+00 |
| 25% | -7.851180e-01 |
| 50% | -2.558327e-01 |
| 75% | 6.704165e-01 |
| max | 2.721397e+00 |

```
In [54]: recurrence_cat_oneHot.columns
```



```
Out[54]: Index(['Thyroid Function', 'Physical Examination', 'Pathology', 'Focality',
              'Risk', 'T', 'N', 'M', 'Stage', 'Response', 'Adenopathy'],
              dtype='object')
```

```
In [55]: recurrence_cat_ordinal.columns
```

```
Out[55]: Index(['Gender', 'Smoking', 'Hx Smoking', 'Hx Radiothreapy'], dtype='object')
```

```
In [56]: recurrence_tar_ordinal.columns
```

```
Out[56]: Index(['Recurred'], dtype='object')
```

```
In [57]: recurrence_num.columns
```

```
Out[57]: Index(['Age'], dtype='object')
```

```
In [58]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder


num_features = ['Age']
cat_features_1hot = ['Thyroid Function', 'Physical Examination', 'Adenopathy', 'Pathol
                  'Risk', 'T', 'N', 'M', 'Stage', 'Response']
cat_features_ordinal = ['Gender', 'Smoking', 'Hx Smoking', 'Hx Radiothreapy']

OneHot_pipeline = Pipeline([
    ("impute", SimpleImputer(strategy='most_frequent')),
    ('1hot_encoder', OneHotEncoder(sparse_output=False, handle_unknown='ignore'))
])

Ordinal_pipeline = Pipeline([
    ("impute", SimpleImputer(strategy='most_frequent')),
    ('ordinal_encoder', OrdinalEncoder())
])
```

```
In [59]: f_preprocessing = ColumnTransformer([
    ("num", num_pipeline, num_features),
    ("1hot", OneHot_pipeline, cat_features_1hot),
    ("feat_ordinal", Ordinal_pipeline, cat_features_ordinal)
])
```

```
In [60]: f_preprocessing
```



```
In [61]: recurrence.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 383 entries, 0 to 382
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    383 non-null    int64
1   Gender                                383 non-null    object
2   Smoking                               383 non-null    object
3   Hx Smoking                            383 non-null    object
4   Hx Radiothreapy                       383 non-null    object
5   Thyroid Function                      383 non-null    object
6   Physical Examination                 383 non-null    object
7   Adenopathy                           383 non-null    object
8   Pathology                            383 non-null    object
9   Focality                             383 non-null    object
10  Risk                                  383 non-null    object
11  T                                      383 non-null    object
12  N                                      383 non-null    object
13  M                                      383 non-null    object
14  Stage                                 383 non-null    object
15  Response                              383 non-null    object
dtypes: int64(1), object(15)
memory usage: 48.0+ KB
```

```
In [62]: recurrence_prepared = f_preprocessing.fit_transform(recurrence)
```

```
In [63]: recurrence_prepared.head()
```

```
Out[63]:
```

| | num_Age | 1hot_Thyroid Function_Clinical Hyperthyroidism | 1hot_Thyroid Function_Clinical Hypothyroidism | 1hot_Thyroid Function_Euthyroid | 1hot_Thyroid Function_Subclinical Hyperthyroidism | 1hot_Thyroid Function_Subclinical Hypothyroidism |
|---|-----------|--|---|------------------------------------|---|--|
| 0 | -0.917439 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 1 | -0.454315 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 2 | -0.718957 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 3 | 1.398184 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 4 | 1.398184 | 0.0 | 0.0 | 1.0 | 0.0 | |

5 rows × 51 columns

SMOTE Oversampling

```
In [64]: X = recurrence_prepared
         Y = recurrence_labels
```

```
In [65]: from sklearn.model_selection import train_test_split
         # Split data into train and test set

         X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify = Y,
```

```
In [66]: from imblearn.over_sampling import SMOTE
         smote = SMOTE(sampling_strategy="minority")
         X_train_SMOTE, y_train_SMOTE = smote.fit_resample(X_train, Y_train)
```


```
X_train_SMOTE.shape, y_train_SMOTE.shape
```

```
Out[66]: ((440, 51), (440, 1))
```


```
In [67]: import matplotlib.pyplot as plt
#Helper function for data distribution
#Visualize proportion of Terminated Vs. Active
def show_recurrence_distrib(data):
    count = ""
    if isinstance(data, pd.DataFrame):
        count = data['Recurred'].value_counts()
    else:
        count = data.value_counts()

    count.plot(kind="pie", explode= [0, 0.1], figsize=(6,6), autopct = "%1.1f%%", shadow
    plt.ylabel("Recurrence: Yes or No")
    plt.legend(["No", "Yes"])
    plt.show()

#Visualize the proportion of borrowers
show_recurrence_distrib(Y_train)
```



```
In [68]: show_recurrence_distrib(y_train_SMOTE)
```



Supervised Learning

Logistic Regression

Logistic Regression with SMOTE

```
In [52]: from imblearn.pipeline import Pipeline as imbpipeline
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix

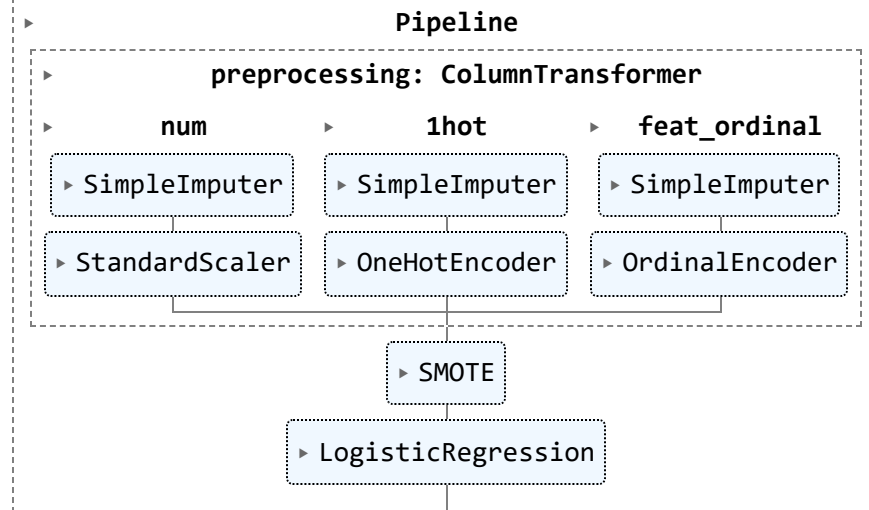
# Define the logistic regression pipeline with SMOTE
logReg_s = imbpipeline(steps=[
    ["preprocessing", f_preprocessing],
    ["SMOTE", SMOTE(random_state=42, sampling_strategy='minority')],
    ["logistic", LogisticRegression(solver='saga', random_state=42, C=1, penalty='l1',
)]

X = recurrence
Y = recurrence_labels
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, r

In [53]: logReg_s.fit(X_train, Y_train)
```

C:\Users\CKY\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
 y = column_or_1d(y, warn=True)

Out[53]:



In [257...

```

y_pred = logReg_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)
P = np.sum(Y_test.values)
N = len(Y_test) - P

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Logistic Regression with SMOTE")
print("-----")
print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

Logistic Regression with SMOTE

```

-----
Precision: 1.0
Accuracy: 0.974025974025974
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.0
F1-score: 0.9523809523809523

```

Logistic Regression with SMOTE and Stratified K-Folds

In [259...

```

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

import warnings
from sklearn.exceptions import DataConversionWarning

# Suppress DataConversionWarning
warnings.filterwarnings(action='ignore', category=DataConversionWarning)

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']
print("Logistic Regression with SMOTE - Stratified K-Folds")
print("-----")
# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    logReg_s.fit(X_train, y_train)

    # Make predictions
    y_pred = logReg_s.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1

```

Logistic Regression with SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.961038961038961
Fold 1 Precision: 1.0
Fold 1 Recall: 0.8636363636363636
Fold 1 F1 Score: 0.9268292682926829

Fold 2 Accuracy: 0.935064935064935
Fold 2 Precision: 0.9473684210526315
Fold 2 Recall: 0.8181818181818182
Fold 2 F1 Score: 0.8780487804878049

Fold 3 Accuracy: 0.948051948051948
Fold 3 Precision: 1.0
Fold 3 Recall: 0.8181818181818182
Fold 3 F1 Score: 0.9

Fold 4 Accuracy: 0.868421052631579
Fold 4 Precision: 0.6896551724137931
Fold 4 Recall: 0.9523809523809523
Fold 4 F1 Score: 0.7999999999999999

Fold 5 Accuracy: 0.7763157894736842
Fold 5 Precision: 0.5526315789473685
Fold 5 Recall: 1.0
Fold 5 F1 Score: 0.711864406779661

```

Logistic Regression with SMOTE - Grid Search

In [261...

```

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

# Define the parameter grid
param_grid = {
    'logistic_C': [0.1, 1, 10], # Regularization parameter
    'logistic_penalty': ['l1', 'l2'] # Penalty term
}

# Perform grid search
grid_search = GridSearchCV(logReg_s, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("Logistic Regression with SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
tn, fp, fn, tp = confusion_matrix_result.ravel()

```

```

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N
f1 = 2 * (precision * recall) / (precision + recall)

# Print metrics
print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

Logistic Regression with SMOTE - Grid Search

```

-----
Best parameters found: {'logistic_C': 1, 'logistic_penalty': 'l1'}
Best F1-score on validation data: 0.9325421396628826
Precision: 1.0
Accuracy: 0.974025974025974
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.0
F1-score: 0.9523809523809523

```

Logistic Regression with SMOTE and PCA

In [263...

```

from sklearn.decomposition import PCA

logReg_s = imbpipeline(steps=[
    ["preprocessing", f_preprocessing],
    ["pca", PCA(n_components=0.95)], # Specify the desired explained variance ratio
    ["SMOTE", SMOTE(random_state=42, sampling_strategy='minority')],
    ["logistic", LogisticRegression(solver='saga', random_state=42, C=1, penalty='l1',
)]

logReg_s.fit(X_train, Y_train)

y_pred = logReg_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)
P = np.sum(Y_test.values)
N = len(Y_test) - P

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Logistic Regression with SMOTE - PCA")
print("-----")

print("Precision:", precision)

```



```
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

Logistic Regression with SMOTE - PCA

```
-----
Precision: 1.0
Accuracy: 0.974025974025974
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.0
F1-score: 0.9523809523809523
```

Logistic Regression without SMOTE

```
In [54]: from imblearn.pipeline import Pipeline as imbpipeline
from sklearn.linear_model import LogisticRegression

# Define the logistic regression pipeline with SMOTE
logReg = imbpipeline(steps=[
    ["preprocessing", f_preprocessing],
    ["logistic", LogisticRegression(solver='saga', random_state=42, C=10, penalty='l2'
)])

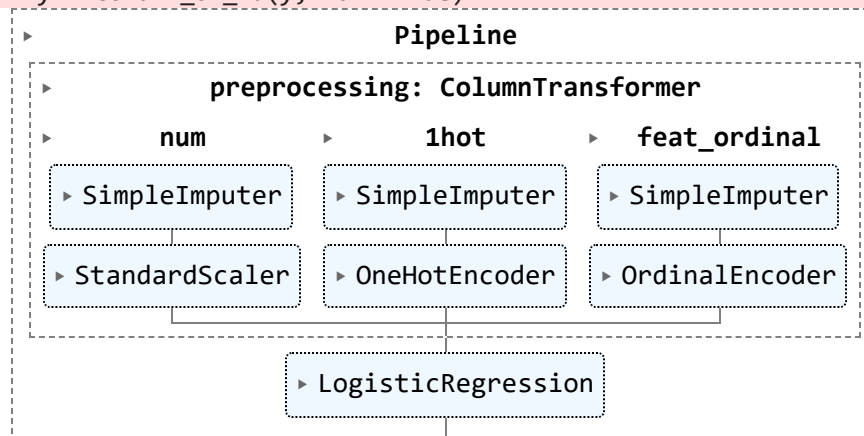
X = recurrence
Y = recurrence_labels
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, r
```

```
In [55]: logReg.fit(X_train, Y_train)
```

C:\Users\CKY\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Out[55]:



```
In [308... y_pred = logReg.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)
P = np.sum(Y_test.values)
N = len(Y_test) - P
```

```

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Logistic Regression without SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

```

Logistic Regression without SMOTE
-----
Precision: 0.9523809523809523
Accuracy: 0.961038961038961
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.01818181818181818
F1-score: 0.9302325581395349

```

Logistic Regression without SMOTE - Stratified K-Folds

In [309...

```

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("Logistic Regression without SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    logReg.fit(X_train, y_train)

    # Make predictions
    y_pred = logReg.predict(X_test)

    # Print evaluation metrics

```

```

print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
print()

```

```

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1

```

Logistic Regression without SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.948051948051948
Fold 1 Precision: 1.0
Fold 1 Recall: 0.8181818181818182
Fold 1 F1 Score: 0.9

Fold 2 Accuracy: 0.922077922077922
Fold 2 Precision: 0.9
Fold 2 Recall: 0.8181818181818182
Fold 2 F1 Score: 0.8571428571428572

Fold 3 Accuracy: 0.948051948051948
Fold 3 Precision: 1.0
Fold 3 Recall: 0.8181818181818182
Fold 3 F1 Score: 0.9

Fold 4 Accuracy: 0.8289473684210527
Fold 4 Precision: 0.6666666666666666
Fold 4 Recall: 0.7619047619047619
Fold 4 F1 Score: 0.7111111111111111

Fold 5 Accuracy: 0.7368421052631579
Fold 5 Precision: 0.5121951219512195
Fold 5 Recall: 1.0
Fold 5 F1 Score: 0.6774193548387097

```

Logistic Regression without SMOTE - Grid Search

In [310...

```

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

# Define the parameter grid
param_grid = {
    'logistic_C': [0.1, 1, 10], # Regularization parameter
    'logistic_penalty': ['l1', 'l2'] # Penalty term
}

# Perform grid search
grid_search = GridSearchCV(logReg, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("Logistic Regression without SMOTE - Grid Search")

```

```

print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N
f1 = 2 * (precision * recall) / (precision + recall)

# Print metrics
print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

Logistic Regression without SMOTE - Grid Search

```

-----
Best parameters found: {'logistic__C': 10, 'logistic__penalty': 'l2'}
Best F1-score on validation data: 0.9404634581105169
Precision: 0.9523809523809523
Accuracy: 0.961038961038961
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.01818181818181818
F1-score: 0.9302325581395349

```

Logistic Regression without SMOTE, with PCA

In [311...

```

from sklearn.decomposition import PCA

logReg = imbpipeline(steps=[
    ["preprocessing", f_preprocessing],
    ["pca", PCA(n_components=0.95)], # Specify the desired explained variance ratio
    ["logistic", LogisticRegression(solver='saga', random_state=42, C=10, penalty='l2'
)])

logReg.fit(X_train, Y_train)

y_pred = logReg.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)
P = np.sum(Y_test.values)
N = len(Y_test) - P

tn, fp, fn, tp = confusion_matrix_result.ravel()

```

```

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Logistic Regression without SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

Logistic Regression without SMOTE - PCA

```

-----
Precision: 0.95
Accuracy: 0.948051948051948
Recall: 0.8636363636363636
True positive rate: 0.8636363636363636
False positive rate: 0.01818181818181818
F1-score: 0.9047619047619048

```

KNN

with SMOTE

```

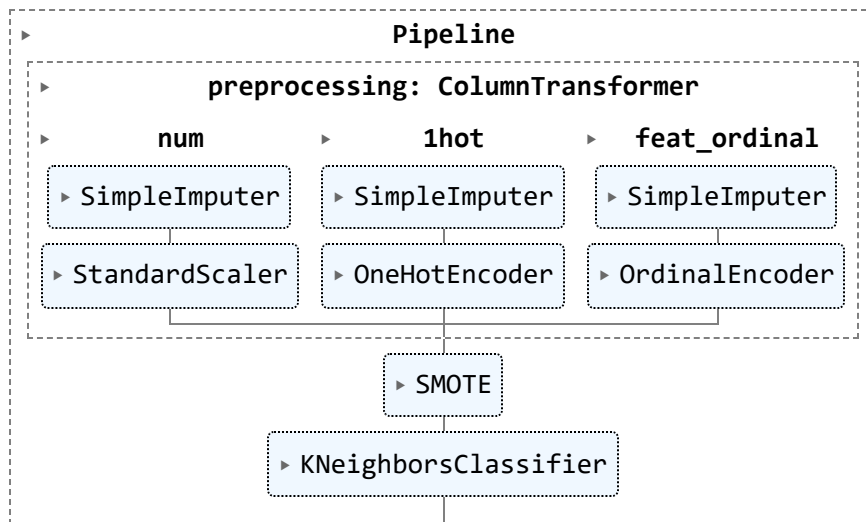
In [56]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix

knn_classifier_s = imbpipeline(steps = [
    ["preprocessing", f_preprocessing],
    ["SMOTE", SMOTE(random_state=42, sampling_strategy='mi
    ['knn', KNeighborsClassifier(n_neighbors=2, weights='u

knn_classifier_s

```

Out[56]:



```

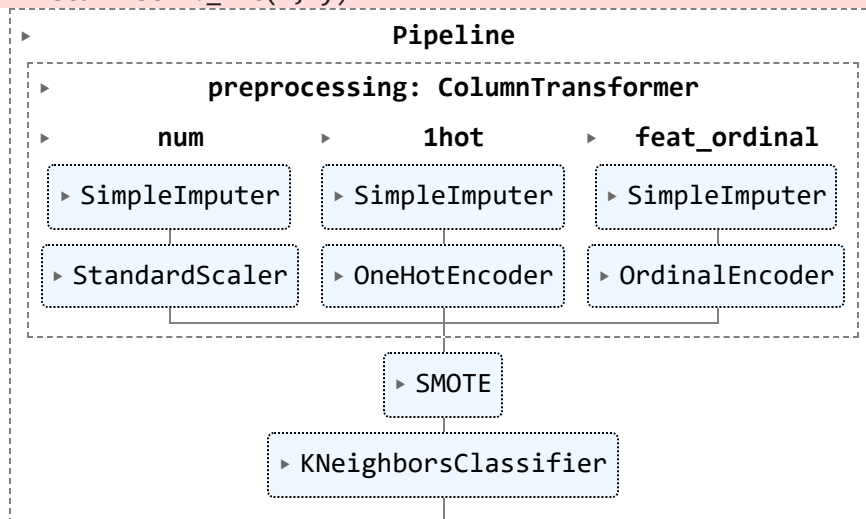
In [57]: X = recurrence
Y = recurrence_labels
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, r
knn_classifier_s.fit(X_train, Y_train)

```

C:\Users\CKY\anaconda3\Lib\site-packages\sklearn\neighbors_classification.py:215: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return self._fit(X, y)

Out[57]:



```

In [347... y_pred = knn_classifier_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)
P = np.sum(Y_test.values)
N = len(Y_test) - P

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

```

```
f1 = 2 * (precision * recall) / (precision + recall)

print("KNN with SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

```
KNN with SMOTE
-----
Precision: 0.9523809523809523
Accuracy: 0.961038961038961
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.01818181818181818
F1-score: 0.9302325581395349
```

KNN with SMOTE - Stratified-K Folds

In [348...

```
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("KNN with SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    knn_classifier_s.fit(X_train, y_train)

    # Make predictions
    y_pred = knn_classifier_s.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
```

```

for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1

```

KNN with SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.8831168831168831
Fold 1 Precision: 1.0
Fold 1 Recall: 0.5909090909090909
Fold 1 F1 Score: 0.7428571428571429

Fold 2 Accuracy: 0.8961038961038961
Fold 2 Precision: 0.8888888888888888
Fold 2 Recall: 0.7272727272727273
Fold 2 F1 Score: 0.7999999999999999

Fold 3 Accuracy: 0.9090909090909091
Fold 3 Precision: 1.0
Fold 3 Recall: 0.6818181818181818
Fold 3 F1 Score: 0.8108108108108109

Fold 4 Accuracy: 0.8421052631578947
Fold 4 Precision: 0.6551724137931034
Fold 4 Recall: 0.9047619047619048
Fold 4 F1 Score: 0.7599999999999999

Fold 5 Accuracy: 0.7105263157894737
Fold 5 Precision: 0.4878048780487805
Fold 5 Recall: 0.9523809523809523
Fold 5 F1 Score: 0.6451612903225806

```

KNN with SMOTE - Grid Search

In [349...

```

from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier

# Define the parameter grid
param_grid = {
    'knn__n_neighbors': [1,2,3,4,5,6,7,8,9], # Test different number of neighbors
    'knn__weights': ['uniform', 'distance'] # Test different weighting schemes
}

# Perform grid search
grid_search = GridSearchCV(knn_classifier_s, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("KNN with SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data

```



```

best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N
f1 = 2 * (precision * recall) / (precision + recall)

# Print metrics
print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

KNN with SMOTE - Grid Search

```

-----
Best parameters found: {'knn__n_neighbors': 2, 'knn__weights': 'uniform'}
Best F1-score on validation data: 0.8582983193277312
Precision: 0.9523809523809523
Accuracy: 0.961038961038961
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.01818181818181818
F1-score: 0.9302325581395349

```

KNN with SMOTE and PCA

In [358...

```

from sklearn.decomposition import PCA

knn_classifier_s = imbpipeline(steps = [
    ["preprocessing", f_preprocessing],
    ["pca", PCA(n_components=0.95)],
    ["SMOTE", SMOTE(random_state=42, sampling_strategy='minority'),
    ["knn", KNeighborsClassifier(n_neighbors=3, weights='uniform')])

knn_classifier_s.fit(X_train, Y_train)

y_pred = knn_classifier_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)
P = np.sum(Y_test.values)
N = len(Y_test) - P

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

```

```

f1 = 2 * (precision * recall) / (precision + recall)

print("KNN with SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

KNN with SMOTE - PCA

```

-----
Precision: 0.875
Accuracy: 0.948051948051948
Recall: 0.9545454545454546
True positive rate: 0.9545454545454546
False positive rate: 0.05454545454545454
F1-score: 0.9130434782608695

```

KNN without SMOTE

```

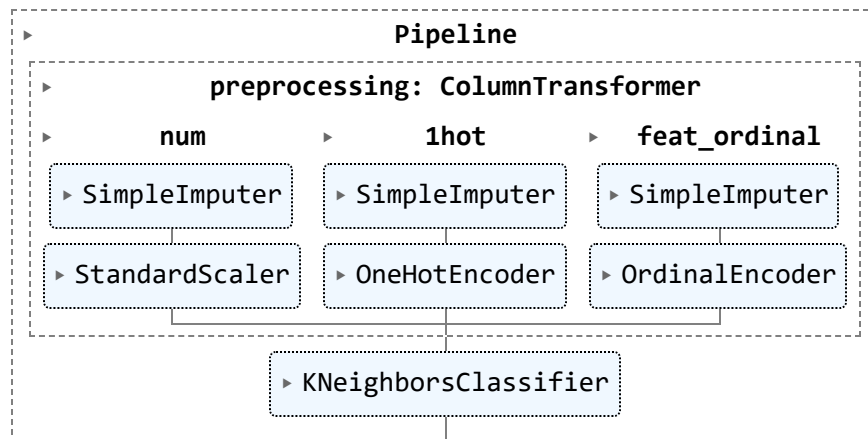
In [58]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix

knn_classifier = imbpipeline(steps = [
                                ["preprocessing", f_preprocessing],
                                ['knn', KNeighborsClassifier(n_neighbors=4, weights='c')])

knn_classifier

```

Out[58]:



```

In [59]: X = recurrence
Y = recurrence_labels
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, r

knn_classifier.fit(X_train, Y_train)

```

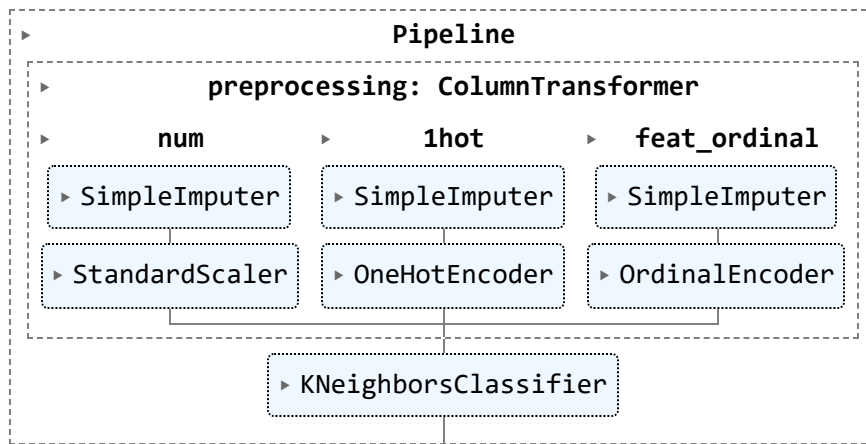
C:\Users\CKY\anaconda3\Lib\site-packages\sklearn\neighbors_classification.py:215: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

return self._fit(X, y)

```

Out[59]:



In [368...

```

y_pred = knn_classifier.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)
P = np.sum(Y_test.values)
N = len(Y_test) - P

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("KNN without SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

KNN without SMOTE

```

-----
Precision: 0.9090909090909091
Accuracy: 0.948051948051948
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.03636363636363636
F1-score: 0.9090909090909091

```

KNN without SMOTE - Stratified-K folds

In [369...

```

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column

```

```
target = df.loc[:, 'Recurred']

print("KNN without SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    knn_classifier.fit(X_train, y_train)

    # Make predictions
    y_pred = knn_classifier.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1
```

KNN without SMOTE - Stratified K-Folds

```
-----
Fold 1 Accuracy: 0.8701298701298701
Fold 1 Precision: 1.0
Fold 1 Recall: 0.5454545454545454
Fold 1 F1 Score: 0.7058823529411764
```

```
Fold 2 Accuracy: 0.8571428571428571
Fold 2 Precision: 0.9230769230769231
Fold 2 Recall: 0.5454545454545454
Fold 2 F1 Score: 0.6857142857142856
```

```
Fold 3 Accuracy: 0.9090909090909091
Fold 3 Precision: 1.0
Fold 3 Recall: 0.6818181818181818
Fold 3 F1 Score: 0.8108108108108109
```

```
Fold 4 Accuracy: 0.8421052631578947
Fold 4 Precision: 0.6551724137931034
Fold 4 Recall: 0.9047619047619048
Fold 4 F1 Score: 0.7599999999999999
```

```
Fold 5 Accuracy: 0.75
Fold 5 Precision: 0.5263157894736842
Fold 5 Recall: 0.9523809523809523
Fold 5 F1 Score: 0.6779661016949152
```

KNN without SMOTE - Grid Search

In [370...

```
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier

# Define the parameter grid
param_grid = {
    'knn_n_neighbors': [1,2,3,4,5,6,7,8,9], # Test different number of neighbors
    'knn_weights': ['uniform', 'distance'] # Test different weighting schemes
}

# Perform grid search
grid_search = GridSearchCV(knn_classifier, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("KNN without SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
tn, fp, fn, tp = confusion_matrix_result.ravel()
```

```

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N
f1 = 2 * (precision * recall) / (precision + recall)

# Print metrics
print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

KNN without SMOTE - Grid Search

```

-----
Best parameters found: {'knn__n_neighbors': 4, 'knn__weights': 'distance'}
Best F1-score on validation data: 0.866474583563833
Precision: 0.9090909090909091
Accuracy: 0.948051948051948
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.03636363636363636
F1-score: 0.9090909090909091

```

KNN without SMOTE, with PCA

In [376...

```

from sklearn.decomposition import PCA

knn_classifier = imbpipeline(steps = [
    ["preprocessing", f_preprocessing],
    ["pca", PCA(n_components=0.95)],
    ["knn", KNeighborsClassifier(n_neighbors=4, weights='c

knn_classifier.fit(X_train, Y_train)

y_pred = knn_classifier.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("KNN without SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)

```

```
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

KNN without SMOTE - PCA

```
-----
Precision: 0.9523809523809523
Accuracy: 0.961038961038961
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.01818181818181818
F1-score: 0.9302325581395349
```

Decision Tree

with SMOTE

```
In [419... from sklearn.tree import DecisionTreeClassifier

# Define your pipeline
decision_tree_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['decision_tree', DecisionTreeClassifier(criterion='entropy', max_depth=10, min_s
```

```
In [420... X = recurrence
Y = recurrence_labels
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, r
```

```
In [421... # Train the model
decision_tree_s.fit(X_train, Y_train)
# Predictions on the test set
y_pred = decision_tree_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Decision Tree with SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

Decision Tree with SMOTE

```
-----
Precision: 1.0
Accuracy: 0.961038961038961
Recall: 0.8636363636363636
True positive rate: 0.8636363636363636
False positive rate: 0.0
F1-score: 0.9268292682926829
```

Decision Tree with SMOTE - Stratified K Folds

In [422...

```
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("Decision Tree with SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    decision_tree_s.fit(X_train, y_train)

    # Make predictions
    y_pred = decision_tree_s.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1
```


Decision Tree with SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.961038961038961
Fold 1 Precision: 1.0
Fold 1 Recall: 0.8636363636363636
Fold 1 F1 Score: 0.9268292682926829

Fold 2 Accuracy: 0.922077922077922
Fold 2 Precision: 0.9444444444444444
Fold 2 Recall: 0.7727272727272727
Fold 2 F1 Score: 0.85

Fold 3 Accuracy: 0.935064935064935
Fold 3 Precision: 0.9473684210526315
Fold 3 Recall: 0.8181818181818182
Fold 3 F1 Score: 0.8780487804878049

Fold 4 Accuracy: 0.8421052631578947
Fold 4 Precision: 0.6551724137931034
Fold 4 Recall: 0.9047619047619048
Fold 4 F1 Score: 0.7599999999999999

Fold 5 Accuracy: 0.7105263157894737
Fold 5 Precision: 0.48717948717948717
Fold 5 Recall: 0.9047619047619048
Fold 5 F1 Score: 0.6333333333333333

```

Decision Tree with SMOTE - Grid Search

In [423...

```

from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'decision_tree__max_depth': [None, 10, 20],
    'decision_tree__min_samples_split': [2, 5, 10],
    'decision_tree__min_samples_leaf': [1, 2, 4]
}

# Perform grid search
grid_search = GridSearchCV(decision_tree_s, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("Decision Tree with SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
# Calculate TP, FP, TN, FN, precision, accuracy, recall, true positive rate, false pos
# Print metrics

```

Decision Tree with SMOTE - Grid Search

Best parameters found: {'decision_tree__max_depth': 20, 'decision_tree__min_samples_leaf': 4, 'decision_tree__min_samples_split': 10}

Best F1-score on validation data: 0.8969971205265324

Decision Tree with SMOTE and PCA

In [425...

```
from sklearn.decomposition import PCA

decision_tree_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ["pca", PCA(n_components=0.95)],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['classifier', DecisionTreeClassifier(criterion='entropy', max_depth=10, min_sampl
]])

decision_tree_s.fit(X_train, Y_train)

y_pred = decision_tree_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Decision Tree with SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

Decision Tree with SMOTE - PCA

Precision: 0.7916666666666666

Accuracy: 0.8961038961038961

Recall: 0.8636363636363636

True positive rate: 0.8636363636363636

False positive rate: 0.09090909090909091

F1-score: 0.8260869565217391

Decision Tree without SMOTE

In [459...

```
from sklearn.tree import DecisionTreeClassifier

# Define your pipeline
decision_tree = imbpipeline([
```

```
['preprocessing', f_preprocessing],
['classifier', DecisionTreeClassifier(criterion='entropy', min_samples_leaf=1, mir
)]
```

```
In [460... X = recurrence
Y = recurrence_labels
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, r
```

```
In [465... # Train the model
decision_tree.fit(X_train, Y_train)
# Predictions on the test set
y_pred = decision_tree.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Decision Tree without SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

```
Decision Tree without SMOTE
-----
Precision: 0.95
Accuracy: 0.948051948051948
Recall: 0.8636363636363636
True positive rate: 0.8636363636363636
False positive rate: 0.01818181818181818
F1-score: 0.9047619047619048
```

Decision Tree without SMOTE - Stratified K Folds

```
In [466... from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("Decision Tree without SMOTE - Stratified K-Folds")
print("-----")
```

```

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    decision_tree.fit(X_train, y_train)

    # Make predictions
    y_pred = decision_tree.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1

```

Decision Tree without SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.961038961038961
Fold 1 Precision: 1.0
Fold 1 Recall: 0.8636363636363636
Fold 1 F1 Score: 0.9268292682926829

Fold 2 Accuracy: 0.935064935064935
Fold 2 Precision: 0.9473684210526315
Fold 2 Recall: 0.8181818181818182
Fold 2 F1 Score: 0.8780487804878049

Fold 3 Accuracy: 0.935064935064935
Fold 3 Precision: 0.9473684210526315
Fold 3 Recall: 0.8181818181818182
Fold 3 F1 Score: 0.8780487804878049

Fold 4 Accuracy: 0.9473684210526315
Fold 4 Precision: 0.9047619047619048
Fold 4 Recall: 0.9047619047619048
Fold 4 F1 Score: 0.9047619047619048

Fold 5 Accuracy: 0.7368421052631579
Fold 5 Precision: 0.5128205128205128
Fold 5 Recall: 0.9523809523809523
Fold 5 F1 Score: 0.6666666666666666

```

Decision Tree without SMOTE - Grid Search

In [475...

```

from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'classifier__max_depth': [None, 10, 20],
    'classifier__min_samples_split': [2, 5, 10],
    'classifier__min_samples_leaf': [1, 2, 4]
}

# Perform grid search
grid_search = GridSearchCV(decision_tree, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("Decision Tree without SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
# Calculate TP, FP, TN, FN, precision, accuracy, recall, true positive rate, false pos
# Print metrics

```

Decision Tree without SMOTE - Grid Search

Best parameters found: {'classifier__max_depth': 20, 'classifier__min_samples_leaf': 4, 'classifier__min_samples_split': 10}

Best F1-score on validation data: 0.8189851616322205

Decision Tree without SMOTE, with PCA

In [473...

```

from sklearn.decomposition import PCA

decision_tree = imbpipeline([
    ['preprocessing', f_preprocessing],
    ["pca", PCA(n_components=0.98)],
    ['classifier', DecisionTreeClassifier(criterion='entropy', min_samples_leaf=1, mir
]])

decision_tree.fit(X_train, Y_train)

y_pred = decision_tree.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

```

```
f1 = 2 * (precision * recall) / (precision + recall)

print("Decision Tree without SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

```
Decision Tree without SMOTE - PCA
-----
Precision: 0.782608695652174
Accuracy: 0.8831168831168831
Recall: 0.8181818181818182
True positive rate: 0.8181818181818182
False positive rate: 0.09090909090909091
F1-score: 0.8
```

Regression Tree

with SMOTE

In [477...

```
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

# Define your pipeline
regression_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['regressor', DecisionTreeRegressor()] # Regression tree
])

# Train the model
regression_s.fit(X_train, Y_train)

y_pred = regression_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Regression Tree with SMOTE")
print("-----")
```

```
print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

Regression Tree with SMOTE

```
-----
Precision: 0.9523809523809523
Accuracy: 0.961038961038961
Recall: 0.9090909090909091
True positive rate: 0.9090909090909091
False positive rate: 0.01818181818181818
F1-score: 0.9302325581395349
```

Regression Tree with SMOTE - Stratified K Folds

In [483...

```
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("Regression Tree with SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    regression_s.fit(X_train, y_train)

    # Make predictions
    y_pred = regression_s.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1
```

Regression Tree with SMOTE - Stratified K-Folds

```
-----
Fold 1 Accuracy: 0.8441558441558441
Fold 1 Precision: 1.0
Fold 1 Recall: 0.45454545454545453
Fold 1 F1 Score: 0.625
```

```
Fold 2 Accuracy: 0.8311688311688312
Fold 2 Precision: 0.7142857142857143
Fold 2 Recall: 0.6818181818181818
Fold 2 F1 Score: 0.6976744186046512
```

```
Fold 3 Accuracy: 0.8961038961038961
Fold 3 Precision: 0.8888888888888888
Fold 3 Recall: 0.7272727272727273
Fold 3 F1 Score: 0.7999999999999999
```

```
Fold 4 Accuracy: 0.7105263157894737
Fold 4 Precision: 0.4838709677419355
Fold 4 Recall: 0.7142857142857143
Fold 4 F1 Score: 0.5769230769230769
```

```
Fold 5 Accuracy: 0.7763157894736842
Fold 5 Precision: 0.5555555555555556
Fold 5 Recall: 0.9523809523809523
Fold 5 F1 Score: 0.7017543859649122
```

Regression Tree with SMOTE - Grid Search

In [485...

```
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
from sklearn.tree import DecisionTreeRegressor

# Define the parameter grid
param_grid = {
    'regressor__max_depth': [None, 10, 20],
    'regressor__min_samples_split': [2, 5, 10],
    'regressor__min_samples_leaf': [1, 2, 4]
}

# Perform grid search
grid_search = GridSearchCV(regression_s, param_grid, cv=5, scoring='neg_mean_squared_e
grid_search.fit(X_train, Y_train)

print("Regression Tree with SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best negative mean squared error on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute evaluation metrics
mse = mean_squared_error(Y_test, y_pred)
print("Mean Squared Error on test data:", mse)
```


Regression Tree with SMOTE - Grid Search

Best parameters found: {'regressor__max_depth': 10, 'regressor__min_samples_leaf': 4, 'regressor__min_samples_split': 10}

Best negative mean squared error on validation data: -0.10672136360468938

Mean Squared Error on test data: 0.06757976591309925

Regression Tree with SMOTE and PCA

In [479...

```

from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

# Define your pipeline
regression_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ["pca", PCA(n_components=0.98)],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['regressor', DecisionTreeRegressor()] # Regression tree
])

# Train the model
regression_s.fit(X_train, Y_train)

y_pred = regression_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Regression Tree with SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

Regression Tree with SMOTE - PCA

Precision: 0.8095238095238095

Accuracy: 0.8831168831168831

Recall: 0.7727272727272727

True positive rate: 0.7727272727272727

False positive rate: 0.07272727272727272

F1-score: 0.7906976744186046

without SMOTE

In [186...

```

from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

# Define your pipeline
regression = imbpipeline([
    ['preprocessing', f_preprocessing],
    ['regressor', DecisionTreeRegressor()] # Regression tree
])

# Train the model
regression.fit(X_train, Y_train)

y_pred = regression.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

```

Precision: 0.9473684210526315
Accuracy: 0.935064935064935
Recall: 0.8181818181818182
True positive rate: 0.8181818181818182
False positive rate: 0.01818181818181818
F1-score: 0.8780487804878049

```

without SMOTE, with PCA

In [187...

```

from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

# Define your pipeline
regression = imbpipeline([
    ['preprocessing', f_preprocessing],
    ["pca", PCA(n_components=0.95)],
    ['regressor', DecisionTreeRegressor()] # Regression tree
])

# Train the model
regression.fit(X_train, Y_train)

y_pred = regression.predict(X_test)

```

```

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

```

Precision: 0.85
Accuracy: 0.8961038961038961
Recall: 0.7727272727272727
True positive rate: 0.7727272727272727
False positive rate: 0.05454545454545454
F1-score: 0.8095238095238095

```

Naive Bayes

with SMOTE

```

In [533... from sklearn.naive_bayes import BernoulliNB

# Create a pipeline with preprocessing, SMOTE, scaling (if necessary), and Bernoulli Naive Bayes
naive_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['naive_bayes', BernoulliNB(alpha=0.1)] # Bernoulli Naive Bayes classifier
])

# Fit the model
naive_s.fit(X_train, Y_train)

# Predictions on the test set
y_pred = naive_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

```

```

print("Naive Bayes with SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

```

Naive Bayes with SMOTE
-----
Precision: 0.9411764705882353
Accuracy: 0.948051948051948
Recall: 0.8421052631578947
True positive rate: 0.7272727272727273
False positive rate: 0.01818181818181818
F1-score: 0.8888888888888888

```

Naive Bayes with SMOTE - Stratified K Folds

In [519...

```

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("Naive Bayes with SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    naive_s.fit(X_train, y_train)

    # Make predictions
    y_pred = naive_s.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):

```

```
# Use the indices provided by split function to extract the corresponding
# train data & test data
train_model(train_index, test_index, fold_no)
fold_no += 1
```

Naive Bayes with SMOTE - Stratified K-Folds

```
-----
Fold 1 Accuracy: 0.8701298701298701
Fold 1 Precision: 0.9285714285714286
Fold 1 Recall: 0.5909090909090909
Fold 1 F1 Score: 0.7222222222222223

Fold 2 Accuracy: 0.9090909090909091
Fold 2 Precision: 1.0
Fold 2 Recall: 0.6818181818181818
Fold 2 F1 Score: 0.8108108108108109

Fold 3 Accuracy: 0.974025974025974
Fold 3 Precision: 1.0
Fold 3 Recall: 0.9090909090909091
Fold 3 F1 Score: 0.9523809523809523

Fold 4 Accuracy: 0.8289473684210527
Fold 4 Precision: 0.625
Fold 4 Recall: 0.9523809523809523
Fold 4 F1 Score: 0.7547169811320755

Fold 5 Accuracy: 0.7236842105263158
Fold 5 Precision: 0.5
Fold 5 Recall: 1.0
Fold 5 F1 Score: 0.6666666666666666
```

Naive Bayes with SMOTE - Grid Search

In [520...

```
from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'naive_bayes__alpha': [0.1, 0.5, 1.0],    # Smoothing parameter
    'naive_bayes__binarize': [0.0, 0.5, 1.0]   # Binarization threshold
}

# Perform grid search
grid_search = GridSearchCV(naive_s, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("Naive Bayes with SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
```

```
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
# Calculate TP, FP, TN, FN, precision, accuracy, recall, true positive rate, false pos
# Print metrics
```

Naive Bayes with SMOTE - Grid Search

Best parameters found: {'naive_bayes__alpha': 0.1, 'naive_bayes__binarize': 0.0}

Best F1-score on validation data: 0.8285456885456887

Naive Bayes with SMOTE - PCA

In [527...

```
from sklearn.decomposition import PCA

naive_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ["pca", PCA(n_components=0.98)],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['naive_bayes', BernoulliNB(alpha=0.1)] # Bernoulli Naive Bayes classifier
])

naive_s.fit(X_train, Y_train)

y_pred = naive_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Naive Bayes with SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

Naive Bayes with SMOTE - PCA

Precision: 0.7272727272727273

Accuracy: 0.8831168831168831

Recall: 0.8421052631578947

True positive rate: 0.7272727272727273

False positive rate: 0.10909090909090909

F1-score: 0.7804878048780488

Naive Bayes without SMOTE

In [522...

```

from sklearn.naive_bayes import BernoulliNB

# Create a pipeline with preprocessing, SMOTE, scaling (if necessary), and Bernoulli Naive Bayes classifier
naive = imbpipeline([
    ['preprocessing', f_preprocessing],
    ['naive_bayes', BernoulliNB(alpha=0.1)] # Bernoulli Naive Bayes classifier
])

# Fit the model
naive.fit(X_train, Y_train)

# Predictions on the test set
y_pred = naive.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("Naive Bayes without SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

```

Naive Bayes without SMOTE
-----
Precision: 0.9411764705882353
Accuracy: 0.948051948051948
Recall: 0.8421052631578947
True positive rate: 0.7272727272727273
False positive rate: 0.01818181818181818
F1-score: 0.8888888888888888

```

Naive Bayes without SMOTE - Stratified K Folds

In [523...

```

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("Naive Bayes without SMOTE - Stratified K-Folds")
print("-----")

```

```

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    naive.fit(X_train, y_train)

    # Make predictions
    y_pred = naive.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1

```

Naive Bayes without SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.8701298701298701
Fold 1 Precision: 0.9285714285714286
Fold 1 Recall: 0.5909090909090909
Fold 1 F1 Score: 0.7222222222222223

Fold 2 Accuracy: 0.9090909090909091
Fold 2 Precision: 1.0
Fold 2 Recall: 0.6818181818181818
Fold 2 F1 Score: 0.8108108108108109

Fold 3 Accuracy: 0.974025974025974
Fold 3 Precision: 1.0
Fold 3 Recall: 0.9090909090909091
Fold 3 F1 Score: 0.9523809523809523

Fold 4 Accuracy: 0.8289473684210527
Fold 4 Precision: 0.625
Fold 4 Recall: 0.9523809523809523
Fold 4 F1 Score: 0.7547169811320755

Fold 5 Accuracy: 0.6973684210526315
Fold 5 Precision: 0.4772727272727273
Fold 5 Recall: 1.0
Fold 5 F1 Score: 0.6461538461538462

```

Naive Bayes without SMOTE -Grid Search

In [526...

```

from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'naive_bayes__alpha': [0.1, 0.5, 1.0],    # Smoothing parameter
    'naive_bayes__binarize': [0.0, 0.5, 1.0]  # Binarization threshold
}

# Perform grid search
grid_search = GridSearchCV(naive, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("Naive Bayes without SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
# Calculate TP, FP, TN, FN, precision, accuracy, recall, true positive rate, false pos
# Print metrics

```

Naive Bayes without SMOTE - Grid Search

 Best parameters found: {'naive_bayes__alpha': 0.1, 'naive_bayes__binarize': 0.0}
 Best F1-score on validation data: 0.8223035327686491

Naive Bayes without SMOTE - PCA

In [531...

```

from sklearn.decomposition import PCA

naive = imbpipeline([
    ['preprocessing', f_preprocessing],
    ['pca', PCA(n_components=0.98)],
    ['naive_bayes', BernoulliNB(alpha=0.1)]
])

naive_s.fit(X_train, Y_train)

y_pred = naive_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

```

```
f1 = 2 * (precision * recall) / (precision + recall)

print("Naive Bayes without SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

```
Naive Bayes without SMOTE - PCA
-----
Precision: 0.7272727272727273
Accuracy: 0.8831168831168831
Recall: 0.8421052631578947
True positive rate: 0.7272727272727273
False positive rate: 0.10909090909090909
F1-score: 0.7804878048780488
```

LDA

with SMOTE

```
In [534... from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

lda_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['lda', LinearDiscriminantAnalysis()] # LDA classifier
])

# Train the model on the SMOTE training data
lda_s.fit(X_train, Y_train)

# Predictions on the test set
y_pred = lda_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("LDA with SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
```

```
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

LDA with SMOTE

```
-----
Precision: 1.0
Accuracy: 0.974025974025974
Recall: 0.8947368421052632
True positive rate: 0.7727272727272727
False positive rate: 0.0
F1-score: 0.9444444444444444
```

LDA with SMOTE - Stratified K Folds

In [536...

```
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("LDA with SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels
    X_train = X.loc[train]
    y_train = y.loc[train]
    X_test = X.loc[test]
    y_test = y.loc[test]

    # Train the model
    lda_s.fit(X_train, y_train)

    # Make predictions
    y_pred = lda_s.predict(X_test)

    # Print evaluation metrics
    print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
    print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
    print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
    print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1
```

LDA with SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.961038961038961
Fold 1 Precision: 1.0
Fold 1 Recall: 0.8636363636363636
Fold 1 F1 Score: 0.9268292682926829

Fold 2 Accuracy: 0.935064935064935
Fold 2 Precision: 0.9473684210526315
Fold 2 Recall: 0.8181818181818182
Fold 2 F1 Score: 0.8780487804878049

Fold 3 Accuracy: 0.8961038961038961
Fold 3 Precision: 0.8181818181818182
Fold 3 Recall: 0.8181818181818182
Fold 3 F1 Score: 0.8181818181818182

Fold 4 Accuracy: 0.7631578947368421
Fold 4 Precision: 0.6153846153846154
Fold 4 Recall: 0.38095238095238093
Fold 4 F1 Score: 0.47058823529411764

Fold 5 Accuracy: 0.8157894736842105
Fold 5 Precision: 0.6
Fold 5 Recall: 1.0
Fold 5 F1 Score: 0.7499999999999999

```

LDA with SMOTE - Grid Search

In [545...

```

from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'lda_n_components': [None, 1, min(X_train.shape[1], len(np.unique(Y_train))) - 1]
}

# Perform grid search
grid_search = GridSearchCV(lda_s, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("LDA with SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
# Calculate TP, FP, TN, FN, precision, accuracy, recall, true positive rate, false pos
# Print metrics

```

LDA with SMOTE - Grid Search

 Best parameters found: {'lda__n_components': None}
 Best F1-score on validation data: 0.8911976911976913

LDA with SMOTE and PCA

In [544...

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

lda_s = imbpipeline([
    ['preprocessing', f_preprocessing],
    ["pca", PCA(n_components=0.95)],
    ['smote', SMOTE(random_state=42, sampling_strategy='minority')],
    ['lda', LinearDiscriminantAnalysis()] # LDA classifier
])

# Train the model on the SMOTE training data
lda_s.fit(X_train, Y_train)

# Predictions on the test set
y_pred = lda_s.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("LDA with SMOTE - PCA")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

LDA with SMOTE - PCA

 Precision: 1.0
 Accuracy: 0.974025974025974
 Recall: 0.8947368421052632
 True positive rate: 0.7727272727272727
 False positive rate: 0.0
 F1-score: 0.9444444444444444

LDA without SMOTE

In [546...

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

lda = imbpipeline([
```

```

        ['preprocessing', f_preprocessing],
        ['lda', LinearDiscriminantAnalysis()] # LDA classifier
    ])

# Train the model on the SMOTE training data
lda.fit(X_train, Y_train)

# Predictions on the test set
y_pred = lda.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("LDA without SMOTE")
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)

```

```

LDA without SMOTE
-----
Precision: 1.0
Accuracy: 0.961038961038961
Recall: 0.8421052631578947
True positive rate: 0.7272727272727273
False positive rate: 0.0
F1-score: 0.9142857142857143

```

LDA without SMOTE - Stratified K Folds

In [548...

```

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

skf = StratifiedKFold(n_splits=5)

# Tell the function which column we are going to use as the target
# Use loc() function to extract the data of that column
target = df.loc[:, 'Recurred']

print("LDA without SMOTE - Stratified K-Folds")
print("-----")

# Define a function to train the model and evaluate on each fold
def train_model(train, test, fold_no):
    X = recurrence
    y = recurrence_labels

```

```

X_train = X.loc[train]
y_train = y.loc[train]
X_test = X.loc[test]
y_test = y.loc[test]

# Train the model
lda.fit(X_train, y_train)

# Make predictions
y_pred = lda.predict(X_test)

# Print evaluation metrics
print('Fold', str(fold_no), 'Accuracy:', accuracy_score(y_test, y_pred))
print('Fold', str(fold_no), 'Precision:', precision_score(y_test, y_pred))
print('Fold', str(fold_no), 'Recall:', recall_score(y_test, y_pred))
print('Fold', str(fold_no), 'F1 Score:', f1_score(y_test, y_pred))
print()

# Perform cross-validation
fold_no = 1
for train_index, test_index in skf.split(df, target):
    # Use the indices provided by split function to extract the corresponding
    # train data & test data
    train_model(train_index, test_index, fold_no)
    fold_no += 1

```

LDA without SMOTE - Stratified K-Folds

```

-----
Fold 1 Accuracy: 0.961038961038961
Fold 1 Precision: 1.0
Fold 1 Recall: 0.8636363636363636
Fold 1 F1 Score: 0.9268292682926829

Fold 2 Accuracy: 0.948051948051948
Fold 2 Precision: 1.0
Fold 2 Recall: 0.8181818181818182
Fold 2 F1 Score: 0.9

Fold 3 Accuracy: 0.922077922077922
Fold 3 Precision: 1.0
Fold 3 Recall: 0.7272727272727273
Fold 3 F1 Score: 0.8421052631578948

Fold 4 Accuracy: 0.9605263157894737
Fold 4 Precision: 0.9090909090909091
Fold 4 Recall: 0.9523809523809523
Fold 4 F1 Score: 0.9302325581395349

Fold 5 Accuracy: 0.75
Fold 5 Precision: 0.525
Fold 5 Recall: 1.0
Fold 5 F1 Score: 0.6885245901639345

```

LDA without SMOTE - Grid Search

In [551...

```

from sklearn.model_selection import GridSearchCV

# Define the parameter grid

```

```

param_grid = {
    'lda__n_components': [None, 1, min(X_train.shape[1], len(np.unique(Y_train))) - 1]
}

# Perform grid search
grid_search = GridSearchCV(lda, param_grid, cv=5, scoring='f1')
grid_search.fit(X_train, Y_train)

print("LDA without SMOTE - Grid Search")
print("-----")

# Print best parameters and best score
print("Best parameters found:", grid_search.best_params_)
print("Best F1-score on validation data:", grid_search.best_score_)

# Evaluate the best model on test data
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# Compute confusion matrix and other metrics
confusion_matrix_result = confusion_matrix(Y_test, y_pred)
# Calculate TP, FP, TN, FN, precision, accuracy, recall, true positive rate, false pos
# Print metrics

```

LDA without SMOTE - Grid Search

Best parameters found: {'lda__n_components': None}

Best F1-score on validation data: 0.9090467276710161

LDA without SMOTE, with PCA

In [552...

```

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

lda = imbpipeline([
    ['preprocessing', f_preprocessing],
    ["pca", PCA(n_components=0.95)],
    ['lda', LinearDiscriminantAnalysis()] # LDA classifier
])

# Train the model on the SMOTE training data
lda.fit(X_train, Y_train)

# Predictions on the test set
y_pred = lda.predict(X_test)

confusion_matrix_result = confusion_matrix(Y_test, y_pred)

tn, fp, fn, tp = confusion_matrix_result.ravel()

precision = tp / (tp + fp)
accuracy = (tp + tn) / (tp + tn + fp + fn)
recall = tp / (tp + fn)
true_positive_rate = tp / P
false_positive_rate = fp / N

f1 = 2 * (precision * recall) / (precision + recall)

print("LDA without SMOTE - PCA")

```



```
print("-----")

print("Precision:", precision)
print("Accuracy:", accuracy)
print("Recall:", recall)
print("True positive rate:", true_positive_rate)
print("False positive rate:", false_positive_rate)
print("F1-score:", f1)
```

LDA without SMOTE - PCA

```
-----
Precision: 1.0
Accuracy: 0.974025974025974
Recall: 0.8947368421052632
True positive rate: 0.7727272727272727
False positive rate: 0.0
F1-score: 0.9444444444444444
```

Unsupervised Learning

Hierachical Clustering

In [195...

```
from sklearn.datasets import make_blobs
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from scipy.cluster.hierarchy import linkage, dendrogram
import matplotlib.pyplot as plt

# Generate synthetic data
X, _ = make_blobs(n_samples=100, centers=5, random_state=42)

# Define the pipeline (although not necessary for hierarchical clustering)
pipeline = Pipeline([
    ('scaler', StandardScaler()), # Scale features if necessary
])

# Fit the pipeline
X1 = pipeline.fit_transform(X)


# Perform hierarchical clustering
Z1 = linkage(X1, method='single', metric='euclidean')
#Z2 = linkage(X1, method='complete', metric='euclidean')
#Z3 = linkage(X1, method='average', metric='euclidean')
#Z4 = linkage(X1, method='ward', metric='euclidean')


plt.figure(figsize=(25, 10))
plt.subplot(2,2,1), dendrogram(Z1), plt.title('Single')
#plt.subplot(3,2,2), dendrogram(Z2), plt.title('Complete')
#plt.subplot(2,2,3), dendrogram(Z3), plt.title('Average')
#plt.subplot(2,2,4), dendrogram(Z4), plt.title('Ward')
plt.show()
```



```
# Plot the clustered data
plt.figure(figsize=(8, 6))
for cluster in range(5):
    plt.scatter(X[cluster_labels == cluster, 0], X[cluster_labels == cluster, 1], s=300)
plt.scatter(pipeline.named_steps['kmeans'].cluster_centers_[0, 0],
            pipeline.named_steps['kmeans'].cluster_centers_[0, 1],
            c='red', s=100, marker='x', label='Cluster Centers')
plt.title('Clustered Data')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.show()
```

C:\Users\CKY\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
 warnings.warn(
C:\Users\CKY\anaconda3\Lib\site-packages\sklearn\cluster_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
 warnings.warn(





In []: